

Car Price Prediction

(1) Import Python Libraries

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import Lasso
from sklearn import metrics
```

(2) Loading the Data Set

```
In [2]: path = r'D:\IITG\portfolio_finance\car_price\car_details.csv'
```

```
In [3]: car_dataset = pd.read_csv(path)
car_dataset.head()
```

```
Out[3]:
```

| | Car_Name | Year | Selling_Price | Present_Price | Kms_Driven | Fuel_Type | Seller_Ty |
|---|----------|------|---------------|---------------|------------|-----------|-----------|
| 0 | ritz | 2014 | 3.35 | 5.59 | 27000 | Petrol | I |
| 1 | sx4 | 2013 | 4.75 | 9.54 | 43000 | Diesel | I |
| 2 | ciaz | 2017 | 7.25 | 9.85 | 6900 | Petrol | I |
| 3 | wagon r | 2011 | 2.85 | 4.15 | 5200 | Petrol | I |
| 4 | swift | 2014 | 4.60 | 6.87 | 42450 | Diesel | I |

```
In [4]: car_dataset.tail()
```

```
Out[4]:
```

| | Car_Name | Year | Selling_Price | Present_Price | Kms_Driven | Fuel_Type | Seller_Ty |
|-----|----------|------|---------------|---------------|------------|-----------|-----------|
| 296 | city | 2016 | 9.50 | 11.6 | 33988 | Diesel | I |
| 297 | brio | 2015 | 4.00 | 5.9 | 60000 | Petrol | I |
| 298 | city | 2009 | 3.35 | 11.0 | 87934 | Petrol | I |
| 299 | city | 2017 | 11.50 | 12.5 | 9000 | Diesel | I |
| 300 | brio | 2016 | 5.30 | 5.9 | 5464 | Petrol | I |

(3) Exploring the Data Set

```
In [5]: car_dataset.shape
```

```
Out[5]: (301, 9)
```

```
In [6]: car_dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 301 entries, 0 to 300
Data columns (total 9 columns):
 #   Column          Non-Null Count  Dtype  
---  -
 0   Car_Name        301 non-null   object 
 1   Year            301 non-null   int64  
 2   Selling_Price   301 non-null   float64 
 3   Present_Price   301 non-null   float64 
 4   Kms_Driven      301 non-null   int64  
 5   Fuel_Type       301 non-null   object 
 6   Seller_Type     301 non-null   object 
 7   Transmission    301 non-null   object 
 8   Owner           301 non-null   int64  
dtypes: float64(2), int64(3), object(4)
memory usage: 21.3+ KB
```

```
In [7]: car_dataset.isnull().sum()
```

```
Out[7]: Car_Name        0
Year                0
Selling_Price       0
Present_Price       0
Kms_Driven          0
Fuel_Type           0
Seller_Type         0
Transmission        0
Owner               0
dtype: int64
```

```
In [8]: car_dataset.Fuel_Type.value_counts()
```

```
Out[8]: Petrol      239
Diesel      60
CNG         2
Name: Fuel_Type, dtype: int64
```

```
In [9]: car_dataset.Seller_Type.value_counts()
```

```
Out[9]: Dealer      195
Individual  106
Name: Seller_Type, dtype: int64
```

```
In [10]: car_dataset.Transmission.value_counts()
```

```
Out[10]: Manual      261
Automatic    40
Name: Transmission, dtype: int64
```

(4) Encoding the Data

```
In [11]: car_dataset.replace({'Fuel_Type':{'Petrol':0,'Diesel':1,'CNG':2}},inplace=True)
car_dataset.replace({'Seller_Type':{'Dealer':0,'Individual':1}},inplace=True)
car_dataset.replace({'Transmission':{'Manual':0,'Automatic':1}},inplace=True)
```

```
In [12]: car_dataset.head()
```

```
Out[12]:
```

| | Car_Name | Year | Selling_Price | Present_Price | Kms_Driven | Fuel_Type | Seller_Ty |
|---|----------|------|---------------|---------------|------------|-----------|-----------|
| 0 | ritz | 2014 | 3.35 | 5.59 | 27000 | 0 | |
| 1 | sx4 | 2013 | 4.75 | 9.54 | 43000 | 1 | |
| 2 | ciaz | 2017 | 7.25 | 9.85 | 6900 | 0 | |
| 3 | wagon r | 2011 | 2.85 | 4.15 | 5200 | 0 | |
| 4 | swift | 2014 | 4.60 | 6.87 | 42450 | 1 | |



(5) Building the Linear Regression Model

Split the Data

```
In [13]: X = car_dataset.drop(['Car_Name', 'Selling_Price'],axis=1)
Y = car_dataset['Selling_Price']
```

```
In [14]: X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size = 0.1, r
```

Create and Fit the Model

```
In [15]: lin_reg_model = LinearRegression()
lin_reg_model.fit(X_train,Y_train)
```

```
Out[15]: LinearRegression()
```

(6) Evaluating the Linear Regression Model

```
In [16]: training_data_prediction = lin_reg_model.predict(X_train)
training_data_prediction[:5]
```

```
Out[16]: array([ 3.73088505,  5.60702081,  7.79779356, -1.88374756,  6.71614572])
```

```
In [19]: test_data_prediction = lin_reg_model.predict(X_test)
test_data_prediction[:5]
```

```
Out[19]: array([10.32892855,  0.77165673,  4.26482324,  4.78985002,  9.88701568])
```

R-squared Error

```
In [17]: error_score = metrics.r2_score(Y_train, training_data_prediction)
print("R squared Error : ", error_score)
```

```
R squared Error :  0.8799451660493701
```

```
In [20]: error_score = metrics.r2_score(Y_test, test_data_prediction)
print("R squared Error : ", error_score)
```


```
R squared Error :  0.836576671502687
```

(7) Visualizing the Linear Regression Model

```
In [18]: plt.scatter(Y_train, training_data_prediction)

plt.xlabel("Actual Price")
plt.ylabel("Predicted Price")
plt.title(" Actual Prices vs Predicted Prices")


plt.show()
```



```
In [21]: plt.scatter(Y_test, test_data_prediction)

plt.xlabel("Actual Price")
plt.ylabel("Predicted Price")
plt.title(" Actual Prices vs Predicted Prices")

plt.show()
```



(8) Building the Lasso Regression Model

Create and Fit the Model

```
In [22]: lass_reg_model = Lasso()
lass_reg_model.fit(X_train,Y_train)
```

Out[22]: Lasso()

(9) Evaluating the Lasso Regression Model

```
In [23]: training_data_prediction = lass_reg_model.predict(X_train)
training_data_prediction[:5]
```

Out[23]: array([3.56679076, 5.60257564, 8.28781371, -0.83081431, 5.2753988])

```
In [24]: test_data_prediction = lass_reg_model.predict(X_test)
test_data_prediction[:5]
```

Out[24]: array([9.87888122, 1.42396266, 4.33267834, 3.17313445, 8.95590579])

R-squared Error

```
In [25]: error_score = metrics.r2_score(Y_train, training_data_prediction)
print("R squared Error : ", error_score)
```

R squared Error : 0.8427856123435794

```
In [26]: error_score = metrics.r2_score(Y_test, test_data_prediction)
print("R squared Error : ", error_score)
```


R squared Error : 0.8709167941173195

(10) Visualizing the Lasso Regression Model

```
In [27]: plt.scatter(Y_train, training_data_prediction)

plt.xlabel("Actual Price")
plt.ylabel("Predicted Price")
plt.title(" Actual Prices vs Predicted Prices")

plt.show()
```



```
In [28]: plt.scatter(Y_test, test_data_prediction)

plt.xlabel("Actual Price")
plt.ylabel("Predicted Price")
plt.title(" Actual Prices vs Predicted Prices")

plt.show()
```

