

# Gold Price Prediction

## (1) Import Python Libraries

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn import metrics
```

## (2) Loading the Data Set

```
In [2]: path = r'D:\IITG\portfolio_finance\gold_price\gold_price_data.csv'
```

```
In [3]: gold_data = pd.read_csv(path)
gold_data.head()
```

```
Out[3]:
```

	Date	SPX	GLD	USO	SLV	EUR/USD
0	1/2/2008	1447.160034	84.860001	78.470001	15.180	1.471692
1	1/3/2008	1447.160034	85.570000	78.370003	15.285	1.474491
2	1/4/2008	1411.630005	85.129997	77.309998	15.167	1.475492
3	1/7/2008	1416.180054	84.769997	75.500000	15.053	1.468299
4	1/8/2008	1390.189941	86.779999	76.059998	15.590	1.557099

```
In [4]: gold_data.tail()
```

```
Out[4]:
```

	Date	SPX	GLD	USO	SLV	EUR/USD
2285	5/8/2018	2671.919922	124.589996	14.0600	15.5100	1.186789
2286	5/9/2018	2697.790039	124.330002	14.3700	15.5300	1.184722
2287	5/10/2018	2723.070068	125.180000	14.4100	15.7400	1.191753
2288	5/14/2018	2730.129883	124.489998	14.3800	15.5600	1.193118
2289	5/16/2018	2725.780029	122.543800	14.4058	15.4542	1.182033

### (3) Exploring the Data Set

In [5]: `gold_data.shape`

Out[5]: (2290, 6)

In [6]: `gold_data.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2290 entries, 0 to 2289
Data columns (total 6 columns):
#   Column      Non-Null Count  Dtype  
---  -
0    Date        2290 non-null   object  
1    SPX         2290 non-null   float64 
2    GLD         2290 non-null   float64 
3    USO         2290 non-null   float64 
4    SLV         2290 non-null   float64 
5    EUR/USD     2290 non-null   float64 
dtypes: float64(5), object(1)
memory usage: 107.5+ KB
```

In [7]: `gold_data.isnull().sum()`

```
Out[7]: Date        0
SPX            0
GLD            0
USO            0
SLV            0
EUR/USD        0
dtype: int64
```

In [8]: `gold_data.describe()`

```
Out[8]:
```


	SPX	GLD	USO	SLV	EUR/USD
<b>count</b>	2290.000000	2290.000000	2290.000000	2290.000000	2290.000000
<b>mean</b>	1654.315776	122.732875	31.842221	20.084997	1.283653
<b>std</b>	519.111540	23.283346	19.523517	7.092566	0.131547
<b>min</b>	676.530029	70.000000	7.960000	8.850000	1.039047
<b>25%</b>	1239.874969	109.725000	14.380000	15.570000	1.171313
<b>50%</b>	1551.434998	120.580002	33.869999	17.268500	1.303297
<b>75%</b>	2073.010070	132.840004	37.827501	22.882500	1.369971
<b>max</b>	2872.870117	184.589996	117.480003	47.259998	1.598798

## (4) Finding Correlations in the Data Set

```
In [9]: correlation = gold_data.corr()
```

```
In [10]: plt.figure(figsize = (8,8))
sns.heatmap(correlation, cbar=True, square=True, fmt='.1f',annot=True, annot
```

```
Out[10]: <AxesSubplot:>
```



```
In [12]: # correlation values of GLD
print(correlation['GLD'])
```


```
SPX      0.049345
GLD      1.000000
USO     -0.186360
SLV      0.866632
EUR/USD  -0.024375
Name: GLD, dtype: float64
```

```
In [16]: # checking the distribution of the GLD Price
sns.distplot(gold_data['GLD'],color='green')
```

C:\Users\SHBHAM\anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

warnings.warn(msg, FutureWarning)

Out[16]: <AxesSubplot:xlabel='GLD', ylabel='Density'>



## (5) Building the Model

### Split the Data

```
In [17]: X = gold_data.drop(['Date', 'GLD'],axis=1)
Y = gold_data['GLD']
```

```
In [18]: X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size = 0.2, r
```

### Create and Fit the Model

```
In [19]: regressor = RandomForestRegressor(n_estimators=100)
regressor.fit(X_train,Y_train)
```


Out[19]: RandomForestRegressor()

## (6) Evaluating the Model

```
In [21]: test_data_prediction = regressor.predict(X_test)
test_data_prediction[:5]
```

```
Out[21]: array([168.77899949,  82.20419984, 115.78319995, 127.67640081,
                120.47140151])
```

### R-squared Error

```
In [22]:  ### R-squared Errorerror_score = metrics.r2_score(Y_test, test_data_prediction)
print("R squared error : ", error_score)
```

```
R squared error :  0.9895529764582724
```

```
In [23]: Y_test = list(Y_test)
```